

# Nearest Neighbors algorithms



Lluís Talavera, 2022



UNIVERSITAT POLITÈCNICA  
DE CATALUNYA  
BARCELONATECH



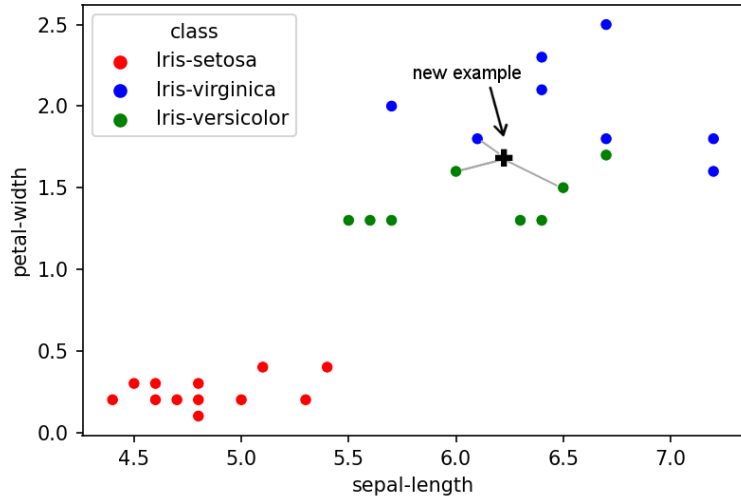
# Nearest neighbors

**Idea:** given a new (unlabeled) example, return the label of the nearest example/s in the training data.

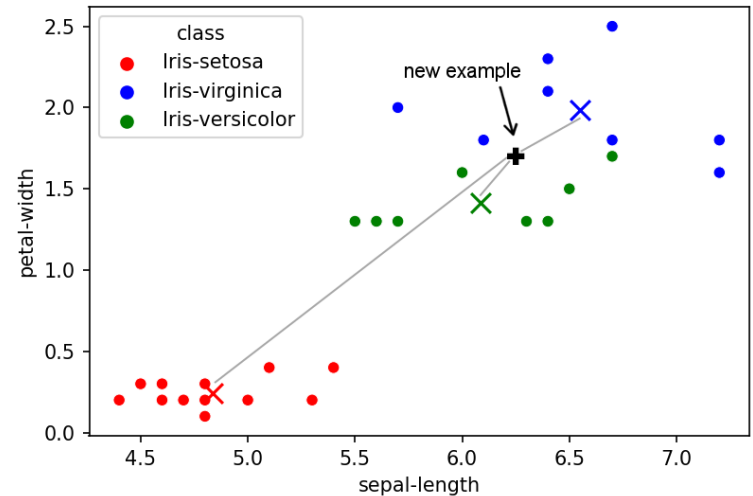
We can:

- summarize examples into centroids: **Nearest Centroid Classifier**
- use the k closest examples: **k-Nearest Neighbor (k-NN)**

3-NN



Nearest Centroid



# Distance metrics

We need a metric to define *nearest*. In the following examples  $n$  is the number of features in the data:

## Euclidean distance

$$d(x,y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

## Manhattan distance

$$d(x,y) = \sum_{i=1}^n |x_i - y_i|$$

## Minkowski distance

$$d(x,y) = (\sum_{i=1}^n (x_i - y_i)^p)^{\frac{1}{p}}$$

$p=1$  for Manhattan distance,  $p=2$  for Euclidean distance

These metrics are for numerical features. There are others for binary and categorical data.

# Nearest Centroid Classifier

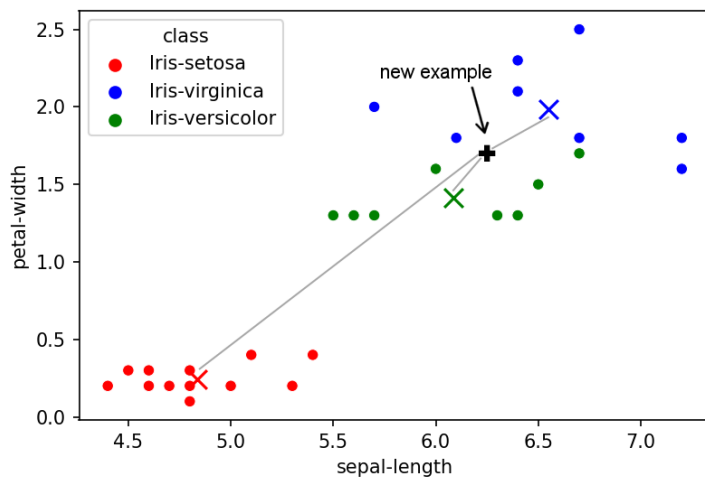
Each class  $C$  is represented by a *centroid*, a kind of "summary example" where the value of each feature  $f_i$  is the average of the values of  $f_i$  across all the examples of  $C$ .

## Training

Compute the centroids for each class.

## Prediction

Compute the distance of the new example to each of the centroids and return the class corresponding to the nearest centroid.



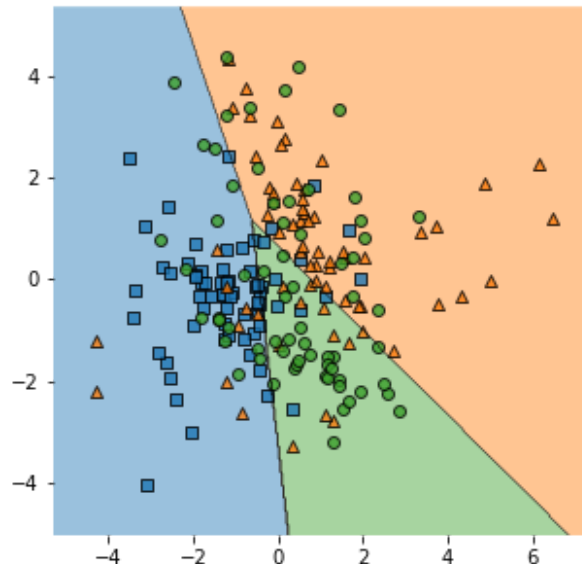
The nearest centroid is from the class *versicolor*

The prediction is *versicolor*

# Nearest Centroid summary

- Too simple but a good choice to start experimenting with classification.
- Could be considered a parametric method (estimate the centroids).
- Performance can degrade for high dimensional data, the distance metric can become distorted.
- Feature with larger values will be dominant in the distance metric (e.g. income over age).
- It is a **linear classifier**.

Decision boundaries:



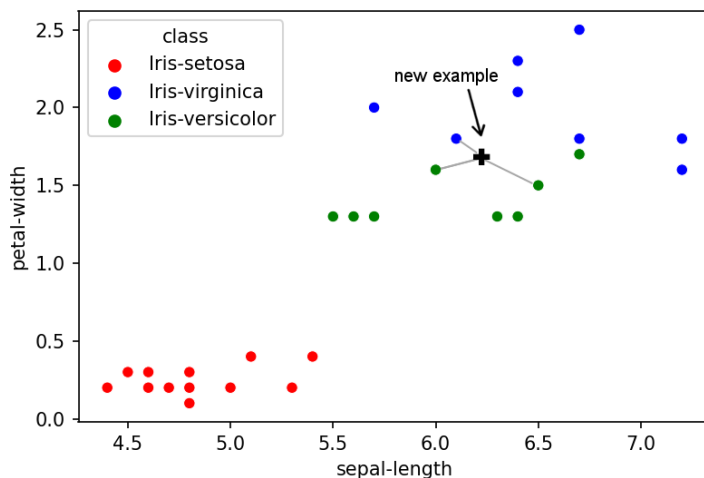
# k-NN for classification

## Training

Just store all the examples (**lazy learning** method)

## Prediction

1. Compute the distance of the new example to each of the stored examples.
2. Select the  $k$  nearest examples (smaller distances).
3. Return the **most frequent class** among the previously selected examples (majority vote).

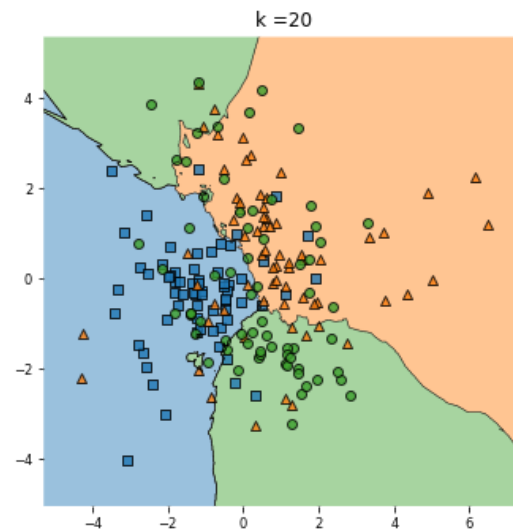
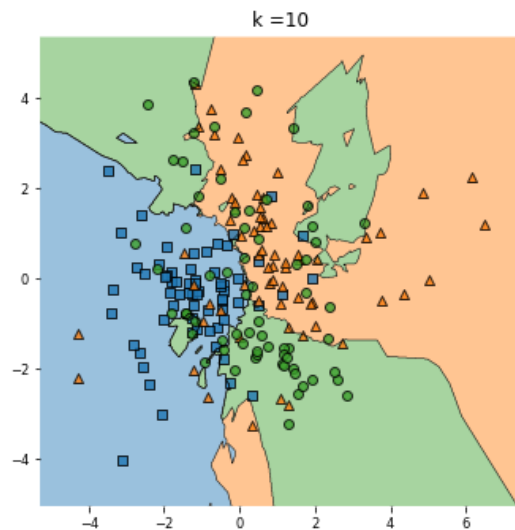
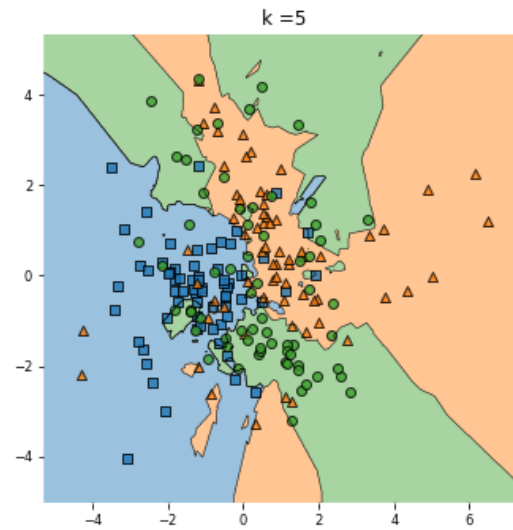
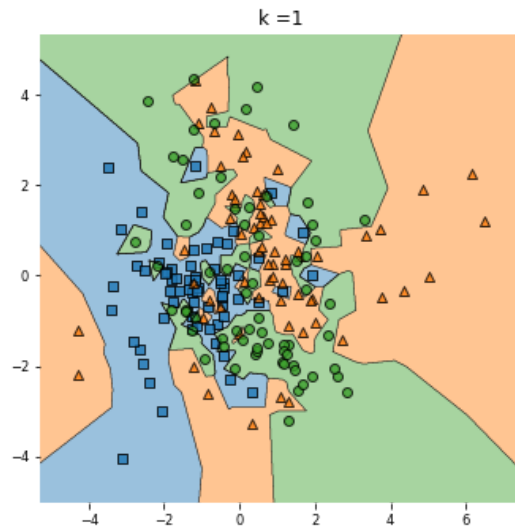


The 3 nearest examples: 2 *versicolor*,  
1 *virginica*

The prediction is *versicolor*

If  $k=1$ , the prediction is *virginica*

# k-NN decision boundaries



# k-NN for regression

Almost the same algorithm but returning a continuous value.

## Training

Just store all the examples

## Prediction

1. Compute the distance of the new example to each of the stored examples.
2. Select the  $k$  nearest examples (smaller distances).
3. Return the **mean of the target** among the previously selected examples.



# k-NN summary

- It is a non-parametric technique.
- Can be used for both classification and regression.
- Common practice is to choose odd values of  $k$  in order to avoid ties.
- The best  $k$  depends on the data
- Performance can degrade for high dimensional data, the distance metric can become distorted.
- Feature with larger values will be dominant in the distance metric (e.g. income over age).
- With large datasets, it can be computationally expensive to retrieve the nearest neighbors. Tree data structures are commonly used to organize the examples and improve retrieval.
- Variant: weight points, close neighbors have greater influence

# imports with sklearn

```
from sklearn.neighbors import NearestCentroid
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.neighbors import KNeighborsRegressor
```

